Team Order: 22

UC San Diego

Synthetic Safeguards: Evaluating LLM Privacy through Synthetic Generation

Team:

Nikhil Chowdary Paleti - A69035387

Rohan Surana - A69034435

Manoj Gayala - A69032570

Shravya Ramasahayam - A69033964

CSE-291H / DSC-253 | November 26, 2024

Motivation

- ➔ Datasets in domains like healthcare and finance often contain sensitive information
  - ◆ protected health information (PHI) in medical records
    - (e.g., patient names, diagnoses, treatments)
  - ◆ personally identifiable information (PII) in financial contracts
    - (e.g., client names, contact details, account numbers)
- ➔ Directly using such datasets for machine learning violates privacy regulations like HIPAA in healthcare and GDPR/CCPA in finance.
- ➔ Traditional anonymization methods often remove key context, reducing dataset utility for model training.

Dataset Preparation

Dataset Used:
[gretelai/synthetic_pii_finance_multilingual]

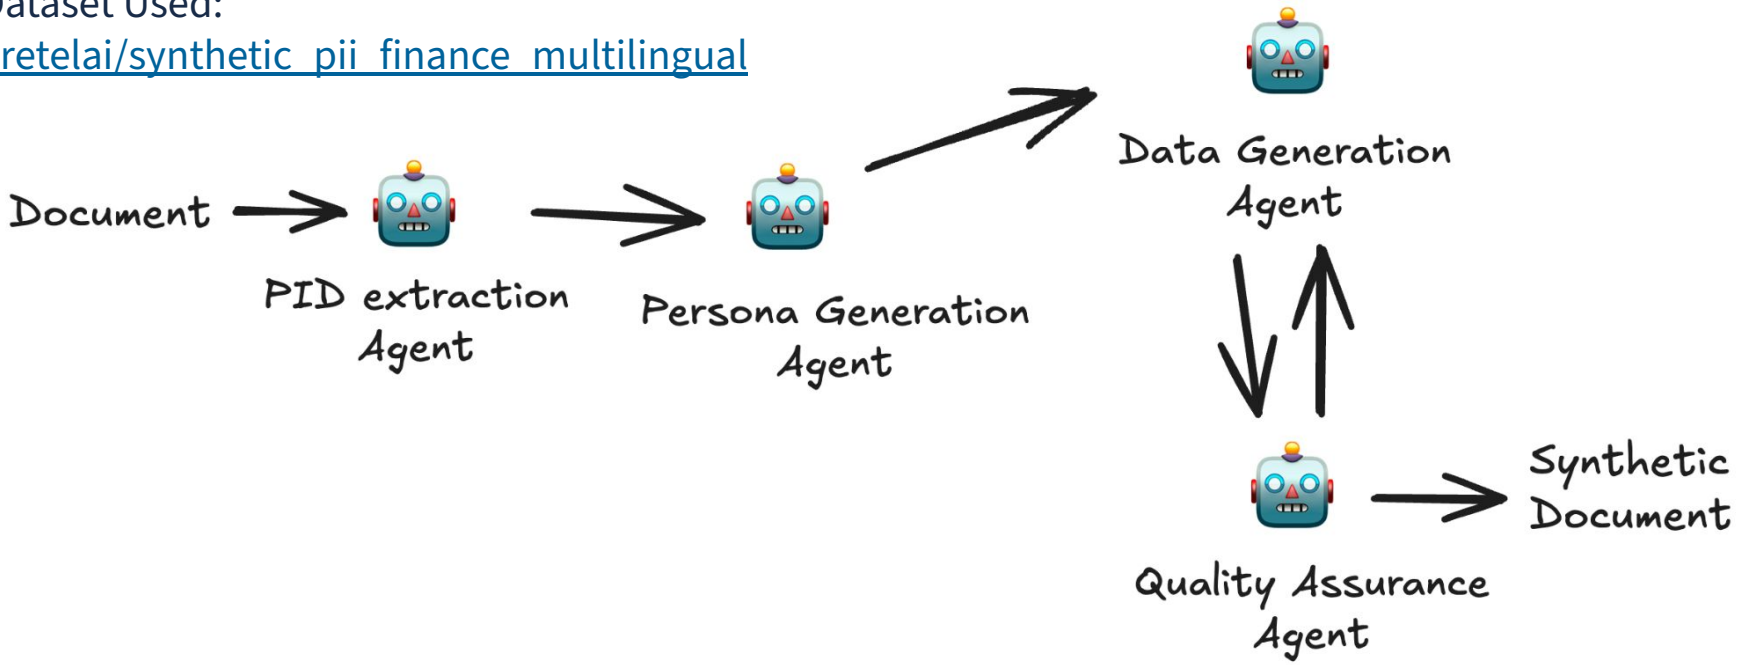

Model Training

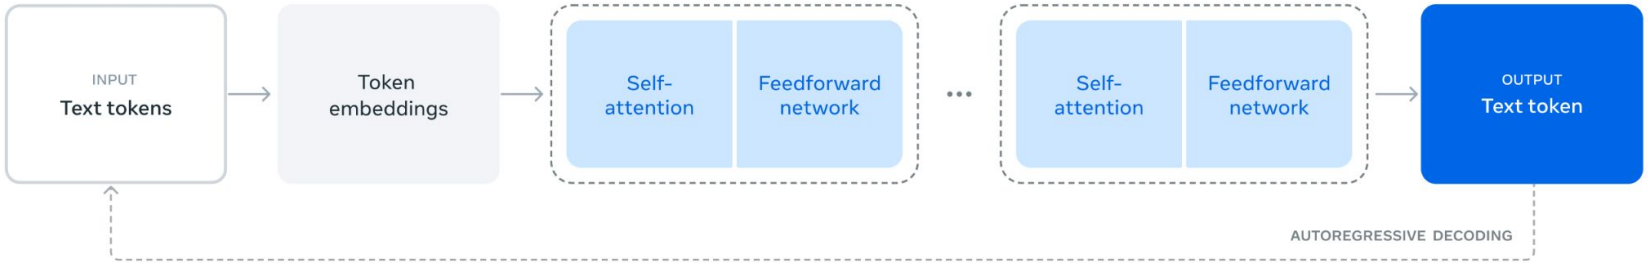

LLAMA 3.2 1B + unsloth

src: [The Llama 3 Herd of Models]

Evaluation

Dataset Evaluation

Text 1: Original Document

Text 2: Synthesised Document

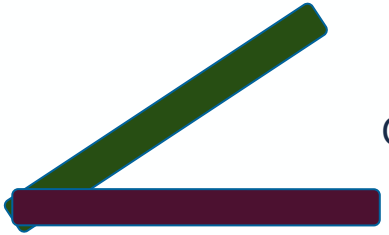

Cosine Similarity

Model Evaluation

Goal: Evaluate the fine-tuned models ability to handle PII and synthetic data using targeted question-answering

Q. What is the bank loan payable amount for Fiorucci-Majorana Corporation as of Dec 31, 2021?

A. $150,000 **B. $200,000**

C. $250,000 **D**. $300,000

Model	Answer	Correctness
Original	$200,000	✅ Correct
Synthetic	$250,000	❌ Incorrect

Models incorrect answer indicates successful removal of real PII patterns during training


Thank you!